





boxes, classes, scores




VISIBLE BACKBONE, YOLO26-N


visible camera



FUSION AT FOUR PYRAMID LEVELS




thermal camera



THERMAL BACKBONE, YOLO26-N, SEPARATE WEIGHTS

PAN NECK WITH STRIDE-4 PATH



HEADS

visible stream thermal stream reliability-weighted sum CMFM: gated cross-modal scan neck detection head

Feature maps are channel-averaged activations of the trained model on this LLVIP test pair; r^θ is its thermal reliability at each fusion level. Shapes are channels $\times$ height² for a 640-pixel input. 6.03 M parameters: 1.37 M per backbone, 2.15 M in the fusers.