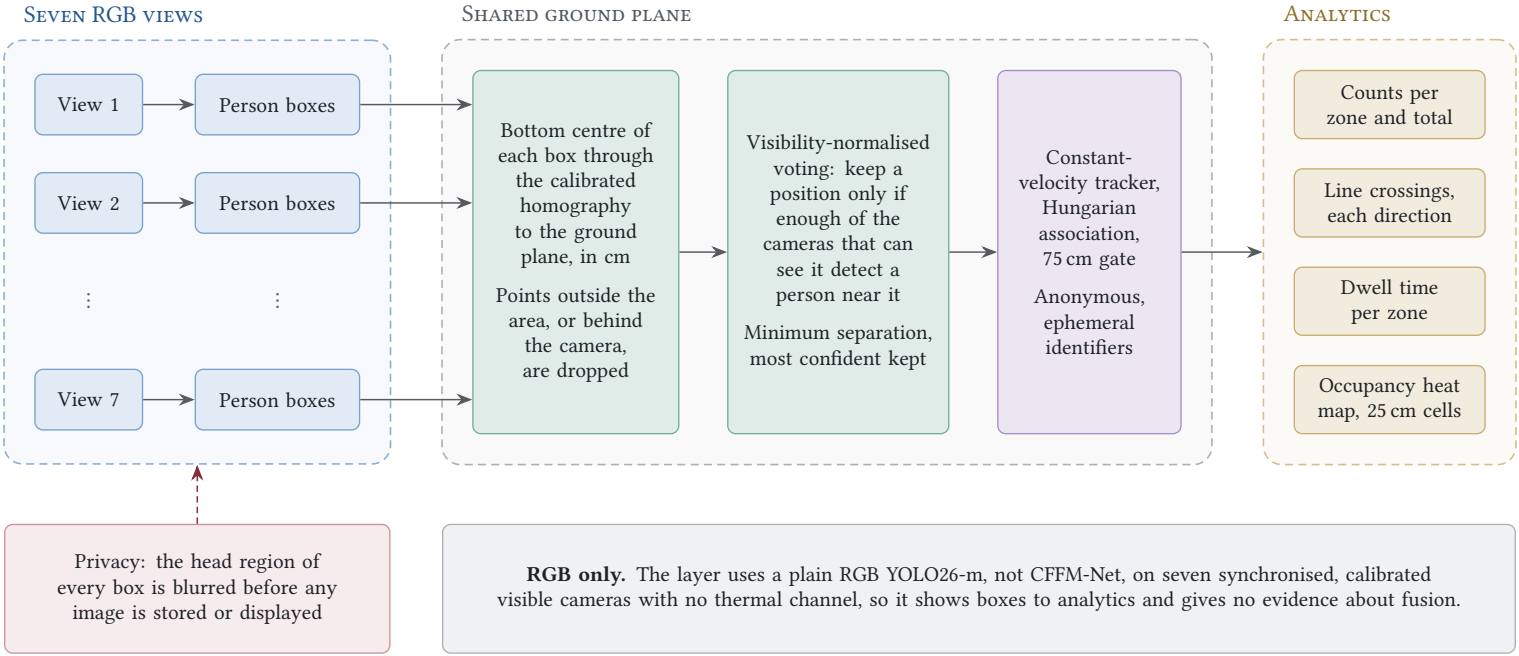
Seven RGB views
View 1
Person boxes
View 2
Person boxes
View 7
Person boxes
Shared ground plane
Bottom centre of each box through the calibrated homography to the ground plane, in cm
Points outside the area, or behind the camera, are dropped
Visibility-normalised voting: keep a position only if enough of the cameras that can see it detect a person near it
Minimum separation, most confident kept
Constant-velocity tracker, Hungarian association, 75 cm gate
Anonymous, ephemeral identifiers
Analytics
Counts per zone and total
Line crossings, each direction
Dwell time per zone
Occupancy heat map, 25 cm cells
Privacy: the head region of every box is blurred before any image is stored or displayed
RGB only. The layer uses a plain RGB YOLO26-m, not CFFM-Net, on seven synchronised, calibrated visible cameras with no thermal channel, so it shows boxes to analytics and gives no evidence about fusion.
WILDTRACK, detector zero-shot and fine-tuned. Detection: precision, recall, MODA, MODP. Tracking: MOTA, IDF1, identity switches, mostly tracked. Counting: mean absolute error, total and per zone; line crossings and dwell per zone